

BIGDATA HADOOP AND SPARK DEVELOPER

STOCK EXCHANGE DATA ANALYSIS PROJECT - SNAPSHOTS

RAVEENA PRABAKARAN

TIGER ANALYTICS

1) Create a data pipeline using sqoop to pull the data from the table below from MYSQL server into Hive.

-> stockcompanies

```
hive> select * from stockcompanies;
OK
COTY Coty, Inc. Consumer Staples Personal Products New York; NY
CEI Crown Castle International Corp. Real Estate REITs Houston; Texas
CSRA CSRA Inc. Information Technology IT Consulting & Other Services Falls Church; Virginia
CSX CSX Corp. Industrials Railroads Jacksonville; Florida
CMI Cummins Inc. Industrials Industrial Machinery Columbus; Indiana
CVS CVS Health Consumer Staples Drug Retail Moonsocket; Rhode Island
DHI D. R. Horton Consumer Discretionary Homebuilding Fort Worth; Texas
DHR Danaher Corp. Industrials Industrial Conglomerates Washington; D.C.
DRI Darden Restaurants Consumer Discretionary Restaurants Orlando; Florida
DVA DaVita Inc. Health Care Health Care Facilities Denver; Colorado
DE Deere & Co. Industrials Construction & Farm Machinery & Heavy Trucks Moline; Illinois
DLPH Delphi Automotive Consumer Discretionary Auto Parts & Equipment Gillingham; Kent; United Kingdom
DAL Delta Air Lines Industrials Airlines Atlanta; Georgia
KRAY Dentsply Sirona Health Care Health Care Supplies York; Pennsylvania
DWN Devon Energy Corp. Energy Oil & Gas Exploration & Production Oklahoma City; Oklahoma
DLR Digital Realty Trust Real Estate Specialized REITs San Francisco; California
DFS Discover Financial Services Financials Consumer Finance Riverwoods; Illinois
DISCA Discovery Communications-A Consumer Discretionary Cable & Satellite Silver Spring; Maryland
DISCK Discovery Communications-C Consumer Discretionary Cable & Satellite Silver Spring; Maryland
DG Dollar General Consumer Discretionary General Merchandise Stores Goodlettsville; Tennessee
DLTR Dollar Tree Consumer Discretionary General Merchandise Stores Chesapeake; Virginia
D Dominion Resources Utilities Electric Utilities Richmond; Virginia
DOV Dow Corpn. Industrials Industrial Machinery Downers Grove; Illinois
DOW Dow Chemical Materials Diversified Chemicals Midland; Michigan
DPS Dr Pepper Snapple Group Consumer Staples Soft Drinks Plano; Texas
DTE DTE Energy Co. Utilities Multiutilities Detroit; Michigan
DD Du Pont (E.I.) Materials Diversified Chemicals Wilmington; Delaware
DUK Duke Energy Utilities Electric Utilities Charlotte; North Carolina
DWB Dun & Bradstreet Industrials Research & Consulting Services Short Hills; New Jersey
ETFC E*Trade Financials Investment Banking & Brokerage New York; New York
EMN Eastman Chemical Materials Diversified Chemicals Kingsport; Tennessee
ETW Eaton Corporation Industrials Industrial Conglomerates Dublin; Ireland
EBAY eBay Inc. Information Technology Internet Software & Services San Jose; California
ECL Ecolab Inc. Materials Specialty Chemicals St. Paul; Minnesota
```

-> stockprices

```
hive> select * from stockprice limit 5;
OK
05-01-2016 WLTW 123.43 125.839996 122.309998 126.25 2163600
06-01-2016 WLTW 125.239998 119.980003 119.940002 125.540001 2386400
07-01-2016 WLTW 116.379997 114.949997 114.93 119.739998 2489500
08-01-2016 WLTW 115.480003 116.620003 113.5 117.440002 2006300
11-01-2016 WLTW 117.010002 114.970001 114.089996 117.330002 1408600
Time taken: 0.067 seconds, Fetched: 5 row(s)
hive>
```

2) Create a new hive table with the following fields by joining the above two hive tables. Please use appropriate Hive built-in functions for columns (a,b,e and h to l).

- Trading_year: Should contain YYYY for each record
- Trading_month: Should contain MM or MMM for each record
- Symbol: Ticker code
- CompanyName: Legal name of the listed company
- State: State to be extracted from headquarters value.
- Sector: Business vertical of the listed company
- Sub_Industry: Business domain of the listed company within a sector
- Open: Average of intra-day opening price by month and year for each listed company
- Close: Average of intra-day closing price by month and year for each listed company
- Low: Average of intra-day lowest price by month and year for each listed company
- High: Average of intra-day highest price by month and year for each listed company
- Volume: Average of number of shares traded by month and year for each listed company

-> stock_data5

```
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 20.81 sec HDFS Read: 52169279 HDFS Write: 2790018 HDFS EC Read: 0 SUCCESS
Stage-Stage-5: Map: 1 Cumulative CPU: 4.67 sec HDFS Read: 2797801 HDFS Write: 4489513 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 25 seconds 480 msec
OK
Time taken: 45.884 seconds
hive> select * from stock_data5 limit 5;
OK
2010 01 A Agilent Technologies Inc California Health Care Health Care Equipment 21.72 21.61 21.4 21.86 4208442.11
2011 01 A Agilent Technologies Inc California Health Care Health Care Equipment 30.29 30.29 29.96 30.65 4496845.0
2012 01 A Agilent Technologies Inc California Health Care Health Care Equipment 28.54 28.78 28.22 29.08 5069975.0
2013 01 A Agilent Technologies Inc California Health Care Health Care Equipment 31.2 31.26 30.97 31.45 4567819.05
2014 01 A Agilent Technologies Inc California Health Care Health Care Equipment 42.01 42.04 41.66 42.36 3494200.0
Time taken: 0.064 seconds, Fetched: 5 row(s)
hive> █
```

DATA ANALYSIS USING HIVE

3) Find the top five companies that are good for investment

```
2022-03-23 10:14:35,415 Stage-3 map = 0%, reduce = 0%
2022-03-23 10:14:41,518 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.82 sec
2022-03-23 10:14:47,613 Stage-3 map = 100%, reduce = 100%, Cumulative CPU 3.4 sec
MapReduce Total cumulative CPU time: 3 seconds 400 msec
Ended Job = job_1640258093152_29384
Launching Job 6 out of 7
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
22/03/23 10:14:48 INFO client.ConfiguredRMFailoverProxyProvider: Failing over to rm81
Starting Job = job_1640258093152_29385, Tracking URL = http://ip-10-0-21-22.ec2.internal:8088/proxy/application_1640258093152_29385/
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cd66.3.2.p0.1605554/lib/hadoop/bin/hadoop job -kill job_1640258093152_29385
Hadoop job information for Stage-4: number of mappers: 1; number of reducers: 1
2022-03-23 10:14:54,686 Stage-4 map = 0%, reduce = 0%
2022-03-23 10:15:00,805 Stage-4 map = 100%, reduce = 0%, Cumulative CPU 1.77 sec
2022-03-23 10:15:05,903 Stage-4 map = 100%, reduce = 100%, Cumulative CPU 3.6 sec
MapReduce Total cumulative CPU time: 3 seconds 600 msec
Ended Job = job_1640258093152_29385
MapReduce Jobs Launched:
Stage-Stage-11: Map: 1 Cumulative CPU: 4.18 sec HDFS Read: 4497123 HDFS Write: 20808 HDFS EC Read: 0 SUCCESS
Stage-Stage-12: Map: 1 Cumulative CPU: 3.71 sec HDFS Read: 4497118 HDFS Write: 21370 HDFS EC Read: 0 SUCCESS
Stage-Stage-9: Map: 1 Cumulative CPU: 3.1 sec HDFS Read: 26730 HDFS Write: 20808 HDFS EC Read: 0 SUCCESS
Stage-Stage-3: Map: 1 Reduce: 1 Cumulative CPU: 3.4 sec HDFS Read: 25661 HDFS Write: 323 HDFS EC Read: 0 SUCCESS
Stage-Stage-4: Map: 1 Reduce: 1 Cumulative CPU: 3.6 sec HDFS Read: 6073 HDFS Write: 342 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 17 seconds 990 msec
OK
Netflix Inc. 1536.0158311345647
Regeneron 1382.2714681440443
Ulta Salon Cosmetics & Fragrance Inc 1174.9378418697165
United Rentals; Inc. 1064.340239912759
Alaska Air Group Inc 878.5555555555554
Time taken: 106.643 seconds, Fetched: 5 row(s)
```

4) Show the best-growing industry by each state, having at least two or more industries mapped.

-> companyanalysis1

```
22/03/23 10:16:01 INFO client.ConfiguredRMFailoverProxyProvider: Failing over to rm11
Starting Job = job_1640258093152_29471, Tracking URL = http://ip-10-0-21-22.ec2.internal:8088/proxy/application_1640258093152_29471/
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cd66.3.2.p0.1605554/lib/hadoop/bin/hadoop job -kill job_1640258093152_29471
Hadoop job information for Stage-10: number of mappers: 1; number of reducers: 0
2022-03-23 10:16:05,153 Stage-10 map = 0%, reduce = 0%
2022-03-23 10:16:11,252 Stage-10 map = 100%, reduce = 0%, Cumulative CPU 3.16 sec
MapReduce Total cumulative CPU time: 3 seconds 100 msec
Ended Job = job_1640258093152_29471
Writing data to directory hdfs://f0m0s0r0c0l0u0s0r0/hive/warehouse/hdapsproject.db/companyanalysis1
MapReduce Jobs Launched:
Stage-Stage-7: Map: 1 Reduce: 1 Cumulative CPU: 6.36 sec HDFS Read: 4098884 HDFS Write: 13238 HDFS EC Read: 0 SUCCESS
Stage-Stage-11: Map: 1 Cumulative CPU: 3.32 sec HDFS Read: 4497242 HDFS Write: 20808 HDFS EC Read: 0 SUCCESS
Stage-Stage-13: Map: 1 Cumulative CPU: 3.47 sec HDFS Read: 4497233 HDFS Write: 21370 HDFS EC Read: 0 SUCCESS
Stage-Stage-10: Map: 1 Cumulative CPU: 3.16 sec HDFS Read: 40844 HDFS Write: 30009 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 100 msec
OK
Time taken: 83.792 seconds
hive> select * from companyanalysis1 limit 5;
OK
Minnesota Industrial Conglomerates 38 Company 112.60817307092305
Georgia Life & Health Insurance 6166 Inc 38-023158802376
Pennsylvania Electrical Components & Equipment APCTEK Inc 192.15081866792448
Texas Integrated Telecommunications Services ATIS Inc 95.25028227324094
Illinois Pharmaceutical 600M 73.4020577451788
```

-> industry_growth

```
hive> select * from industry_growth limit 50;
OK
California      Apparel Retail      261.8019767671436
California      Application Software 120.31201428397625
California      Biotechnology       189.855597859821
California      Health Care Equipment 147.29321666303753
California      Home Entertainment Software 298.3868203665769
California      Internet Software & Services 336.56041800779366
California      MultiUtilities      62.895455776620366
California      Networking Equipment 16.60504085306744
California      REITs              131.08399426163498
California      Retail REITs       107.50814027970536
California      Semiconductor Equipment 147.9653267908011
California      Semiconductors     282.4593566657631
Connecticut     Industrial Conglomerates 75.02785859563353
Florida Hotels; Resorts & Cruise Lines 137.37933900437363
Florida Industrial Conglomerates 177.39424996713996
Florida Internet Software & Services 163.10267039273256
Illinois        Construction & Farm Machinery & Heavy Trucks 73.11111617775501
Illinois        Health Care Equipment 41.94627917684228
Illinois        Industrial Machinery 136.03903698890892
Illinois        Packaged Foods & Meats 55.38509384373717
Illinois        REITs              65.02753772568181
Ireland Industrial Conglomerates 136.6124307695158
Ireland Pharmaceuticals 245.14488964557233
Maryland        Cable & Satellite   85.87024737549106
Massachusetts   Biotechnology       260.2045747013016
Massachusetts   Health Care Equipment 169.22809835926253
Massachusetts   REITs              68.24052887161051
Massachusetts   Semiconductors     285.17414566777273
Michigan         MultiUtilities      143.76404421430016
Minnesota       Packaged Foods & Meats 166.69719818630523
Missouri        Industrial Conglomerates 85.58764014977949
New Jersey      Health Care Equipment 120.88893239332754
New Jersey      Research & Consulting Services 115.46492537344713
```

-> best-growing industry by each state

```
22/03/23 11:30:05 INFO client.ConfiguredRMFailoverProxyProvider: Failing over to rm81
Starting Job = job_1640258093152_29554, Tracking URL = http://ip-10-0-21-22.ec2.internal:8088/proxy/application_1640258093152_29554/
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/lib/hadoop/bin/hadoop job -kill job_1640258093152_29554
Hadoop job information for Stage-4: number of mappers: 1; number of reducers: 0
2022-03-23 11:30:12,035 Stage-4 map = 0%, reduce = 0%
2022-03-23 11:30:19,213 Stage-4 map = 100%, reduce = 0%, Cumulative CPU 3.65 sec
MapReduce Total cumulative CPU time: 3 seconds 650 msec
Ended Job = job_1640258093152_29554
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.43 sec HDFS Read: 10966 HDFS Write: 833 HDFS EC Read: 0 SUCCESS
Stage-Stage-4: Map: 1 Cumulative CPU: 3.65 sec HDFS Read: 7158 HDFS Write: 1409 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 11 seconds 80 msec
OK
California      Internet Software & Services 336.56041800779366
Connecticut     Industrial Conglomerates 75.02785859563353
Florida Industrial Conglomerates 177.39424996713996
Illinois        Industrial Machinery 136.03903698890892
Ireland Pharmaceuticals 245.14488964557233
Maryland        Cable & Satellite   85.87024737549106
Massachusetts   Semiconductors     285.17414566777273
Michigan         MultiUtilities      143.76404421430016
Minnesota       Packaged Foods & Meats 166.69719818630523
Missouri        Industrial Conglomerates 85.58764014977949
New Jersey      Health Care Equipment 120.88893239332754
New York        Diversified Financial Services 244.27936670090116
North Carolina  Apparel; Accessories & Luxury Goods 234.5365478318533
Ohio Banks      171.58635235361513
Oklahoma        Oil & Gas Exploration & Production 129.8777266777514
Pennsylvania    Diversified Chemicals 163.20832113323664
Texas Airlines  571.0256197386555
United Kingdom  Insurance Brokers 181.0504280362911
Virginia        Aerospace & Defense 255.37107674669656
Washington      Internet & Direct Marketing Retail 323.86488316718413
Wisconsin        Electric Utilities 131.94734372014784
Time taken: 43.652 seconds, Fetched: 21 row(s)
hive>
```

5) For each sector find the following.

-> sectorwisegrowth

```
2022-03-23 11:45:30,105 Stage-5 map = 0%, reduce = 0%
2022-03-23 11:45:30,312 Stage-5 map = 100%, reduce = 0%, Cumulative CPU 3.15 sec
MapReduce Total cumulative CPU time: 3 seconds 150 msec
Ended Job = job_1640258093152_29554
Moving data to directory hdfs://nameservice1/user/hive/warehouse/bdhsproject.db/sectorwisegrowth
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.62 sec HDFS Read: 450062 HDFS Write: 3561 HDFS EC Read: 0 SUCCESS
Stage-Stage-4: Map: 1 Reduce: 1 Cumulative CPU: 6.84 sec HDFS Read: 800978 HDFS Write: 3581 HDFS EC Read: 0 SUCCESS
Stage-Stage-5: Map: 1 Cumulative CPU: 3.15 sec HDFS Read: 9536 HDFS Write: 3007 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 17 seconds 610 msec
OK
Time taken: 61.578 seconds
hive> select * from sectorwisegrowth limit 20;
OK
Consumer Discretionary 2010 14.40480730721993
Consumer Discretionary 2011 4.80012658278584
Consumer Discretionary 2012 7.73096129113922
Consumer Discretionary 2013 24.29257290431553
Consumer Discretionary 2014 9.60380523809531
Consumer Discretionary 2015 6.200380052380959
Consumer Discretionary 2016 14.081984761904747
Consumer Staples 2010 4.18229241170440
Consumer Staples 2011 4.15827411764097
Consumer Staples 2012 4.537941176470589
Consumer Staples 2013 8.81092436074278
Consumer Staples 2014 9.32285742857139
Consumer Staples 2015 4.80361004130937
Consumer Staples 2016 3.18305555555569
Energy 2010 7.71806090909091
Energy 2011 -1.594848484848491
Energy 2012 -1.94909090909091
Energy 2013 12.584444444444436
Energy 2014 -6.264444444444443
Energy 2015 10.090444444444443
Time taken: 0.055 seconds, Fetched: 20 row(s)
hive>
```

- Worst year

```
at org.apache.logging.log4j.core.appender.AbstractManager.getAppender(AbstractManager.java:114)
at org.apache.logging.log4j.core.appender.OutputStreamManager.getAppender(OutputStreamManager.java:114)
at org.apache.logging.log4j.core.appender.RandomAccessFileManager.getAppender(RandomAccessFileManager.java:74)
at org.apache.hadoop.hive.q1.log.HushuhofRandomAccessFileAppender.createAppender(HushuhofRandomAccessFileAppender.java:178)
... 35 more

Execution completed successfully
MapReduce task succeeded
Launching Job 2 out of 2
Number of reduce tasks is set to 0 since there's no reduce operator
22/03/23 11:54:50 INFO ClientConfigurationProvider: Falling over to rml
Starting Job = job_1640258093152_29629, Tracking URL = http://ip-10-0-21-22.ec2.internal:8088/proxy/application_1640258093152_29629/
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/11b/hadoop/bin/hadoop job -kill job_1640258093152_29629
Hadoop job information for Stage-4: number of mappers: 1; number of reducers: 0
2022-03-23 11:54:50.172 Stage-4 map = 0%, reduce = 0%
2022-03-23 11:55:03.272 Stage-4 map = 100%, reduce = 0%, Cumulative CPU 2.88 sec
MapReduce Total cumulative CPU time: 2 seconds 880 msec
Finished Job = job_1640258093152_29629
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.07 sec HDFS Read: 10935 HDFS Write: 536 HDFS EC Read: 0 SUCCESS
Stage-Stage-4: Map: 1 Cumulative CPU: 2.88 sec HDFS Read: 6839 HDFS Write: 646 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 880 msec
OK
Consumer Discretionary 2011 4.86012658227804
Consumer Staples 2011 3.1830555555555556
Energy 2011 -10.093644444444444
Financials 2011 -6.6080905117415
Health Care 2011 2.808080870202485
Industrials 2011 -2.6400000000000002
Information Technology 2011 -2.802519045598095
Materials 2011 -3.967001311313135
Real Estate 2011 -4.46340272680078
Telecommunications Services 2011 -2.2939999999999999
Utilities 2011 -6.473027428166
Time taken: 42.643 seconds, fetched: 11 row(s)
hive>
```

- b. Best year

```
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.util.RunJar.run(RunJar.java:313)
at org.apache.hadoop.util.RunJar.main(RunJar.java:227)

Execution completed successfully
MapReduce task succeeded
Launching Job 2 out of 2
Number of reduce tasks is set to 0 since there's no reduce operator
22/03/23 12:00:26 INFO ClientConfigurationProvider: Falling over to rml
Starting Job = job_1640258093152_29645, Tracking URL = http://ip-10-0-21-22.ec2.internal:8088/proxy/application_1640258093152_29645/
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/11b/hadoop/bin/hadoop job -kill job_1640258093152_29645
Hadoop job information for Stage-4: number of mappers: 1; number of reducers: 0
2022-03-23 12:00:33.731 Stage-4 map = 0%, reduce = 0%
2022-03-23 12:00:39.628 Stage-4 map = 100%, reduce = 0%, Cumulative CPU 3.12 sec
MapReduce Total cumulative CPU time: 3 seconds 120 msec
Finished Job = job_1640258093152_29645
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 8.33 sec HDFS Read: 11059 HDFS Write: 536 HDFS EC Read: 0 SUCCESS
Stage-Stage-4: Map: 1 Cumulative CPU: 3.12 sec HDFS Read: 6943 HDFS Write: 646 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 11 seconds 450 msec
OK
Consumer Discretionary 2013 24.39257258011653
Consumer Staples 2014 9.322857142857139
Energy 2016 18.64222222222223
Financials 2016 16.3266258032924815
Health Care 2014 24.325254237288992
Industrials 2016 19.88492548434434
Information Technology 2013 15.532024249999997
Materials 2016 19.036400000000015
Real Estate 2014 18.892412793103472
Telecommunications Services 2014 5.4999999999999995
Utilities 2014 10.268571428571427
Time taken: 55.214 seconds, fetched: 11 row(s)
hive>
```

- c. Stable year

```
at org.apache.hadoop.hive.q1.exec.mr.MapReduceTask.executeInProcess(MapReduceTask.java:193)
at org.apache.hadoop.hive.q1.exec.mr.ExecDriver.main(ExecDriver.java:779)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.util.RunJar.run(RunJar.java:313)
at org.apache.hadoop.util.RunJar.main(RunJar.java:227)
Caused by: java.lang.IllegalStateException: ManagerFactory [org.apache.logging.log4j.core.appender.RandomAccessFileManager$RandomAccessFileManagerFactory@611f8128]
unable to create manager for [UserJob/JobExecution_Tags/00186247028-4306-a931-e421813b3a/raennaprhuhartigeranaly_20220323120539_ec2ec2-806c-46ec-3e11-c
5c4-94a402d] with data [org.apache.logging.log4j.core.appender.RandomAccessFileManager$FactoryData@79c81d17]
at org.apache.logging.log4j.core.appender.AbstractManager.getAppender(AbstractManager.java:114)
at org.apache.logging.log4j.core.appender.OutputStreamManager.getAppender(OutputStreamManager.java:114)
at org.apache.logging.log4j.core.appender.RandomAccessFileManager.getAppender(RandomAccessFileManager.java:74)
at org.apache.hadoop.hive.q1.log.HushuhofRandomAccessFileAppender.createAppender(HushuhofRandomAccessFileAppender.java:178)
... 42 more

Execution completed successfully
MapReduce task succeeded
Launching Job 2 out of 2
Number of reduce tasks is set to 0 since there's no reduce operator
22/03/23 12:06:08 INFO ClientConfigurationProvider: Falling over to rml
Starting Job = job_1640258093152_29649, Tracking URL = http://ip-10-0-21-22.ec2.internal:8088/proxy/application_1640258093152_29649/
Kill Command = /opt/cloudera/parcels/CDH-6.3.2-1.cdh6.3.2.p0.1605554/11b/hadoop/bin/hadoop job -kill job_1640258093152_29649
Hadoop job information for Stage-4: number of mappers: 1; number of reducers: 0
2022-03-23 12:06:15.321 Stage-4 map = 0%, reduce = 0%
2022-03-23 12:06:22.442 Stage-4 map = 100%, reduce = 0%, Cumulative CPU 3.1 sec
MapReduce Total cumulative CPU time: 3 seconds 180 msec
Finished Job = job_1640258093152_29649
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.24 sec HDFS Read: 11806 HDFS Write: 536 HDFS EC Read: 0 SUCCESS
Stage-Stage-4: Map: 1 Cumulative CPU: 3.1 sec HDFS Read: 6868 HDFS Write: 87 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 360 msec
OK
Time taken: 43.742 seconds
hive>
```